

Lecture Notes on CS231N

Yiwei Chen



Zhejiang University

April 4, 2025

Lecture 1: History of CV and Introduction to CNNs

- **ImageNet:** Annual competition for image classification, started in 2010.
- **Convolutional Neural Networks (CNNs):** Introduced by Yann LeCun in 1998, CNNs are a type of neural network designed for processing structured grid data, such as images. CNNs show great performance in image classification tasks.

Lecture 2: Image Classification Pipeline

1. Attempts:

- Find edges, then corners: does not work well.
- Use large datasets with labels.

2. Classifiers:

(a) K-Nearest Neighbors (KNN):

- *Description:* When $K = 1$, Find the closest image in the dataset to the input image (Nearest Neighbors (NN)).

- *Distance metric:*

- i. L1(Mahanttan) Distance: a squared distance metric.

$$d(x, y) = \sum_i |x_i - y_i| \quad (1)$$

- ii. L2(Euclidean) Distance: a squared distance metric.

$$d(x, y) = \sqrt{\sum_i (x_i - y_i)^2} \quad (2)$$

Rotating the coordinate system changes the L1 distance but not the L2 distance.

- *Performance:*

Training time: $O(1)$, as there is nothing to do.

Prediction time: $O(N)$, which is inefficient.

- K-Nearest Neighbors (KNN):

Description:

When $K = 1$, the classifier is too sensitive to noise.

Instead of copying the label of the closest image, take the majority vote of the K closest images.

- *hyperparameters*(超参数):


Choices about the model that are not learned from the data, e.g., K in KNN.

To set hyperparameters:

- Never use the test set to set hyperparameters.
- Splitting data into train and test is not enough.
- **The better idea:** Splitting the training set into training set, validation set, and test set.


Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data



Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!



- **The common idea:** Cross-validation(交叉验证).

Idea #4: Cross-Validation: Split data into **folds**, try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

- *Pros and Cons:*

Actually, KNN on image is never used:

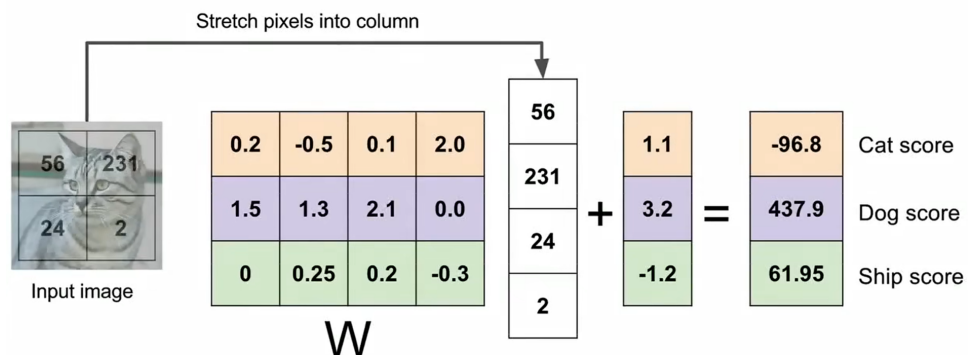
- Very slow at test time.
- Distance-metrics on pixels are not informative.
- Curse of dimensionality: as the number of dimensions increases, the distance between points becomes less meaningful.

(b) **Linear Classifier:**

- *Description:* A linear classifier makes its predictions based on a linear predictor function combining a set of weights with the feature vector.

$$f(x, W) = Wx + b \quad (3)$$

where x is the input image and w is the weight vector, b is the bias term.




- *Hard cases:*

Hard cases for a linear classifier


Class 1:
number of pixels > 0 odd

Class 2:
number of pixels > 0 even



Class 1:
 $1 \leq L2 \text{ norm} \leq 2$

Class 2:
Everything else



Class 1:
Three modes

Class 2:
Everything else

